

Signal Peptide Efficiency Prediction Using Protein Language Model Embeddings and Neural Network Regression

Mehak Wadhwa

Department of Chemistry, Fordham University

Advisor: Dr. Joshua Schrier

June 2026

Abstract

Signal peptides direct protein secretion in bacteria and are critical targets for optimizing recombinant protein production. I compare machine learning approaches for predicting signal peptide efficiency in *Bacillus subtilis*, evaluating random forest and neural network regressors across physicochemical descriptors and several protein language model (PLM) embedding sets. A vector regression approach (predicting 10-bin probability distributions rather than scalar values) combined with multi-seed ensembling gives my reproducible headline: an ensemble of architecturally diverse *open-weight* PLMs (ProtT5, ESM2-650M, and ProtBERT) achieves a test MSE of 0.975 ± 0.003 (mean $\pm$ std over independent retrains), an $\sim 20\%$ improvement over the published random forest baseline of 1.22 [Grasso et al., 2023]. The best single embedding is the proprietary Ginkgo AA-0 (0.965 ± 0.005), but its API has been discontinued and its embeddings can no longer be regenerated by anyone; I therefore report it only as a caveated comparison and anchor all reproducible claims on open weights, recovering essentially all of AA-0’s performance (a linear probe shows AA-0 and ESM2-650M agree at Pearson $r = 0.953$).

Two findings temper these results. A linear probe on the same embeddings captures most of the predictive gain over a mean baseline, indicating that the PLM embedding carries most of the task-relevant signal, with the neural network head adding a smaller further improvement. Leave-one-gene-out cross-validation reveals MSE 2.19 (Spearman $\rho = 0.72$), a sharp increase over the standard split (in which 132 of 134 genes overlap train and test), and is reported as the realistic cross-gene generalization estimate.

Cross-dataset evaluation on four external datasets shows weak but consistent zero-shot transfer in the biologically correct direction (agreement-aligned Spearman ~ 0.2 – 0.3 ; the Wu and Xue random-forest results survive multiple-testing correction); fine-tuning a target-supervised head on frozen features does not help and, for the three continuous datasets, overfits and degrades the transfer, so absolute secretion-efficiency magnitude does not transfer. Applied to 4832 designed variants across 131 genes, the models achieve Spearman $\rho = 0.36$ – 0.39 and classification accuracy of 72–74%, useful for prioritizing candidates in experimental screening.

1 Introduction

Signal peptides are short N-terminal sequences (typically 15–30 amino acids) that direct nascent proteins through cellular secretion machinery. In industrial biotechnology, the choice of signal peptide can determine whether a recombinant protein is produced at commercially viable yields or fails entirely [Grasso et al., 2023]. For organisms like *Bacillus subtilis*, a workhorse of industrial enzyme production, optimizing signal peptide selection is essential for maximizing secretion efficiency.

Traditional approaches to signal peptide optimization rely on experimental screening, which is time-consuming and expensive. Machine learning offers the potential to predict secretion efficiency from sequence alone, enabling rational design of signal peptide–protein combinations. Grasso et al. [2023] constructed a library of signal peptide variants and trained a random forest regressor on 156 physicochemical descriptors, achieving a test MSE of 1.22 for predicting the weighted average (WA) secretion score.

Recent advances in protein language models (PLMs) provide an alternative to hand-crafted features. Models like ESM2 [Lin et al., 2023] learn rich representations of protein sequences from evolutionary data, capturing structural and functional information that may not be encoded in traditional descriptors. However, effectively leveraging these high-dimensional embeddings (1280–2560 dimensions) requires appropriate model architectures and hyperparameter tuning. Existing tools such as SignalP 6.0 [Teufel et al., 2022] and DeepSig [Savojardo et al., 2018] predict signal peptide presence and cleavage sites with high accuracy, but they do not predict quantitative secretion efficiency, which is the focus of this work.

PLM embeddings have shown broad utility for protein property prediction beyond signal peptides. The TAPE benchmark [Rao et al., 2019] demonstrated that pretrained representations transfer effectively to tasks ranging from stability prediction to remote homology detection, establishing PLMs as general-purpose protein feature extractors. ProtTrans [El-naggar et al., 2022] extended this to multiple transformer architectures, showing that models pretrained on billions of protein sequences capture structural and functional signals without explicit evolutionary information. More directly relevant, Meier et al. [2021] showed that ESM-family models can predict mutational effects on protein function in a zero-shot setting, suggesting that PLM representations encode fitness-relevant information even without task-specific fine-tuning. However, most of these studies focus on classification tasks (functional annotation, stability, localization) or per-residue prediction (secondary structure, contact maps). Quantitative regression on continuous functional readouts (particularly for secretion efficiency, where the output is a cell-sorting distribution rather than a single scalar) remains less explored.

In this work, I systematically evaluate signal peptide efficiency prediction across multiple feature representations and model types. I perform hyperparameter search independently for each feature type, use regression rather than classification for continuous efficiency prediction, and scale neural network architectures to the input dimensionality of each embedding.

My analysis reveals a clear interaction between model type and feature representation: neural networks substantially outperform random forests when given PLM embeddings, but underperform on physicochemical features. Beyond this core comparison, I extend the analysis in five directions: (i) a vector regression approach that predicts bin probability dis-

tributions instead of scalar WA, yielding further performance gains; (ii) a reproducible open-weight ensemble of architecturally diverse PLMs (ProtT5, ESM2-650M, ProtBERT) that recovers the performance of the now-discontinued proprietary Ginkgo AA-0 embedding while remaining fully reproducible; (iii) cross-dataset generalization evaluation against four external datasets; (iv) a design task evaluating model utility for ranking novel signal peptide variants; and (v) a leave-one-gene-out analysis that provides a realistic cross-gene generalization estimate.

2 Methods

2.1 Data

All data originate from the supplementary material of Grasso et al. [2023]. The dataset comprises signal peptide variants fused to an α -amylase (AmyQ) reporter in *B. subtilis*, with the weighted average of the per-variant activity-bin distribution used as a proxy for secretion efficiency. Quality filters follow the original protocol: signal peptide length 10–40 amino acids, WA values between 1.0 and 10.0. The original train/test partition is preserved. After applying the 10–40 amino acid length filter, the physicochemical-feature split contains 3085 training and 1323 test samples (13 over-length variants are dropped).

I evaluate four feature representations:

- **PhysChem:** 156 physicochemical descriptors from Grasso et al. [2023], including amino acid composition, hydrophobicity, and secondary structure propensities for each signal peptide region.
- **ESM2-650M:** 1280-dimensional mean-pooled embeddings from the 650M-parameter ESM2 model [Lin et al., 2023].
- **ESM2-3B:** 2560-dimensional embeddings from the 3B-parameter ESM2 model.
- **Ginkgo AA-0:** 1280-dimensional embeddings from Ginkgo Bioworks’ AA-0 protein language model [Ginkgo Bioworks, 2024], trained on proprietary protein sequence data. No peer-reviewed publication describing the model architecture or training data was available at the time of writing. **This model is API-only and has since been discontinued: its weights were never released, and the model API and self-serve registration portal are no longer available.** Its embeddings can therefore no longer be regenerated by anyone (me, the original authors, or a reader), so AA-0 cannot anchor a reproducible result; I report it only as a caveated comparison and base all reproducible results on open-weight models.
- **ProtT5:** 1024-dimensional mean-pooled embeddings from the open-weight ProtT5 model [Elnaggar et al., 2022] (a T5-architecture PLM).
- **ProtBERT:** 1024-dimensional mean-pooled embeddings from the open-weight ProtBERT model [Elnaggar et al., 2022] (a BERT-architecture PLM).

The ProtT5 and ProtBERT embeddings, alongside ESM2-650M, are all open-weight and reproducible from publicly available HuggingFace checkpoints with no API dependency, and together form the open-ensemble headline (Section 3.5).

The PLM datasets retain the full Grasso partition (3095 training and 1326 test samples); the 13-sample difference from the physicochemical split reflects the 10–40 amino acid length filter, which was applied only to the physicochemical pipeline.

2.2 Preprocessing

I standardize features to zero mean and unit variance using statistics computed on the training data only. I evaluate target skewness and apply a $\log(1 + x)$ transformation when $|\text{skew}| > 1.0$. For all datasets in this study, skewness was 0.116, so no transformation was applied.

2.3 Random Forest Models

I first reproduce the baseline result using the exact hyperparameters from Grasso et al. [2023]: 75 trees, max depth 25, `min_samples_split` = 0.001, `min_samples_leaf` = 0.0001, with all features available at each split.

For systematic comparison, I perform randomized hyperparameter search independently for each feature type using 100 configurations evaluated under 5-fold cross-validation. The search space includes: number of trees $\in \{50, 75, 100, 150, 200, 300\}$, max depth $\in \{10, 15, 20, 25, 30, \text{None}\}$, minimum samples for splitting $\in \{2, 5, 10, 0.001, 0.01\}$, minimum samples per leaf $\in \{1, 2, 4, 0.0001, 0.001\}$, and fraction of features per split $\in \{\text{sqrt}, \log2, 0.5, 0.75, 1.0\}$.

2.4 Neural Network Regression

I train feedforward neural networks with the following architecture for each hidden layer: fully connected $\rightarrow$ batch normalization $\rightarrow$ ReLU activation $\rightarrow$ dropout. The output layer is a single neuron with linear activation, trained to minimize MSE.

Hidden layer sizes are scaled to input dimensionality:

- PhysChem (156d): architectures from $\{[128], [128, 64], [256, 128]\}$
- ESM2-650M / Ginkgo AA-0 (1280d): $\{[256], [512, 256], [256, 128]\}$
- ESM2-3B (2560d): $\{[512], [512, 256], [1024, 512]\}$

Additional hyperparameters: dropout rate $\in \{0.2, 0.3, 0.4, 0.5\}$, L2 regularization $\in \{10^{-4}, 10^{-3}, 10^{-2}\}$, learning rate $\in \{10^{-3}, 5 \times 10^{-4}, 10^{-4}\}$, and batch size $\in \{32, 64\}$. I sample 40 configurations per feature type, training each with the Adam optimizer for up to 200 epochs with early stopping (patience 15) and learning rate reduction on plateau. An 80/20 train/validation split is used for model selection; the configuration with the lowest validation MSE is evaluated on the held-out test set.

2.5 Multi-Seed Ensemble Protocol

To reduce variance from random initialization and data splitting, I employ a 5-seed ensemble protocol with seeds {42, 123, 456, 789, 1024}. For each seed, the training set is split 80/20 into training and validation subsets; a model is trained with early stopping on the validation loss. The ensemble prediction is the arithmetic mean of the five per-seed predictions. This protocol is used for the neural network models in Sections 2.6–2.8.

2.6 Vector Regression

Rather than predicting a scalar WA value, I predict the 10-dimensional bin probability distribution from which WA is derived. The WA measurement in Grasso et al. [2023] is computed from fluorescence-activated cell sorting into 10 bins, with $WA = \sum_{i=1}^{10} i \cdot p_i$ where p_i is the fraction of cells in bin i . Following the convention of Grasso et al. [2023], bin 1 contains the highest-secreting cells, so a *lower* WA indicates more efficient secretion. Some sequences exhibit bimodal bin distributions where the WA falls between two peaks (Figure 1), illustrating that the scalar WA can be a poor summary of the underlying data; predicting the full distribution preserves this information.




Figure 1: Bimodal bin probability distribution for gene *aspB*, variant MKLAKRVSAITP-STTLAITQKAKEL (WA = 5.63). The weighted average (dashed line) falls in the valley between peaks at bins 3 and 7, illustrating that the scalar WA can be a poor summary of the underlying distribution and motivating the vector regression approach.

The vector regression architecture uses two hidden layers of 256 units each with LeakyReLU activation and 20% dropout, followed by a 10-unit output layer with softmax activation to produce a valid probability distribution. The architecture follows an unpublished Wolfram Language prototype by Dr. Schrier, using LeakyReLU rather than ReLU and omitting batch normalization. Training uses the Adam optimizer with learning rate 5×10^{-4} , batch size 32, and up to 200 epochs with early stopping (patience 15) and learning rate reduction on plateau. Two loss functions are compared: categorical cross-entropy (CCE) and focal loss

[Lin et al., 2017], defined as $\text{FL}(p_t) = -\alpha(1 - p_t)^\gamma \log(p_t)$ with $\alpha = 0.25$ and $\gamma = 2.0$, where p_t is the model’s predicted probability for bin t . Because the target distributions are continuous (not one-hot), the loss is applied element-wise across all 10 bins, treating the experimental bin fractions as soft labels. Focal loss down-weights well-predicted bins, encouraging the model to focus on samples with ambiguous bin distributions. This is a non-standard application of focal loss, which was originally designed for hard class labels with severe class imbalance in object detection [Lin et al., 2017]; the theoretical justification is weaker for soft probability targets, though the intuition of down-weighting easy predictions still applies. Each embedding–loss combination is trained as a 5-seed ensemble (Section 2.5). After prediction, WA is recovered as $\hat{\text{WA}} = \hat{\mathbf{p}} \cdot [1, 2, \dots, 10]^\top$.

Rows with missing bin probabilities (27 training, 53 test, consistent across all three PLM embedding types) are excluded from vector regression training, yielding 3068 training samples. The validation-split models (Section 3.3) evaluate on the 1273 test samples with complete bin data; the optimized models described below evaluate on all 1326 test samples using WA as the target.

2.7 Architecture and Regularization Optimization

I next explore whether removing the validation split and compensating with explicit regularization can improve upon the base vector regression model. Training on all 3068 samples with a fixed epoch budget of 300 and no early stopping removes the regularization signal provided by validation-based stopping, so I rely instead on dropout tuning and architectural depth to control overfitting. I select the dropout rate on a held-out *validation* split (Appendix E), not on the test set. I explore:

- **Dropout rate:** 0.15, 0.20, 0.25, 0.30, and 0.35, applied to a deeper (256, 256, 128) three-layer architecture with 5-seed ensembles. Dropout is selected on validation data; on validation, all rates in 0.20–0.35 are statistically indistinguishable.
- **Deeper architectures:** 4-layer configurations including (256, 256, 128, 64), (384, 384, 256, 128), and others, with 5-seed ensembles.
- **Cosine annealing:** Learning rate schedule $\eta_t = \eta_0 \cdot \frac{1}{2}(1 + \cos(\pi t/T))$ combined with deeper architectures.
- **Large ensembles:** 20-seed ensembles with the best configurations.
- **Mixed-architecture ensemble:** Averaging predictions from multiple distinct architectures.

All configurations use Ginkgo AA-0 embeddings and focal loss. Each 5-seed ensemble uses seeds {42, 123, 456, 789, 1024} with independent random initialization; predictions are averaged across seeds.

2.8 Cross-Dataset Generalization

Strong within-dataset accuracy does not guarantee practical value; a model is only useful if its predictions transfer to new experimental contexts. To evaluate this, I test the best ESM2-650M RF and the 5-seed NN ensemble on four external datasets:

- **Wu** [Wu et al., 2020]: 81 signal peptides in *E. coli* with binary secretion outcomes (functional/non-functional).
- **Xue** [Xue et al., 2023]: 322 signal peptides in *S. cerevisiae* with enzyme activity measurements.
- **Zhang-P43 / Zhang-PglVM** [Zhang et al., 2016]: 114 signal peptides in *B. subtilis* measured under two different promoters.

Spearman rank correlation serves as the primary metric, as it is invariant to differences in scale and units between datasets. Spearman is invariant to scale but not to *direction*: the external assays score higher values as better secretion, whereas the model predicts Grasso WA where lower is better (bin 1 = best secretor). I therefore orient every cross-dataset correlation so that a positive value indicates agreement (equivalently, I correlate each external metric against the negated predicted WA); rank-correlation magnitudes are unchanged, only the sign. External-metric orientations were confirmed from the data: Xue (enzyme activity, 0–10437) and Zhang (secretion screening, 0–193 / 0–327) are raw magnitudes with higher = better, and Wu uses the standard functional=1 binary encoding. For the Zhang datasets, the same 114 signal peptides are measured under two promoters (P43 and PglVM), which lets me check how well rankings are conserved across promoters.

Transfer learning via fine-tuning. To test whether limited target-domain data strengthens the weak zero-shot transfer, I fine-tune a vector regression ensemble with the same validation-selected architecture (focal loss, (256, 256, 128), validation-selected dropout) trained on ESM2-650M embeddings, which are open-weight and reproducible for external sequences (Ginkgo AA-0 embeddings cannot be generated for external sequences). Fine-tuning uses 5-fold cross-validation with 5 random seeds per fold. I report the *pooled out-of-fold* Spearman correlation (computed once over all held-out predictions concatenated across folds) with a bootstrap 95% confidence interval, rather than the biased mean of per-fold correlations $\pm$ standard deviation, since per-fold correlations on small folds are noisy and their mean is upward-biased in magnitude. Significance for each dataset is judged by whether its bootstrap CI excludes zero; across the four datasets this is not multiple-testing corrected, though the effects I report as significant also pass a four-test Bonferroni threshold. The pretrained model’s first two hidden blocks (6 layers) are frozen; the third hidden block and the 10-unit softmax output head are replaced with a task-specific head: a single sigmoid unit for the binary Wu dataset, or a single linear unit for the continuous datasets (Xue, Zhang-P43, Zhang-PglVM). Features are standardized using statistics from the Grasso training set, ensuring the frozen pretrained layers receive inputs with the same distribution as during pretraining. Fine-tuning uses a reduced learning rate (10^{-4} , $5\times$ lower than pretraining), batch size 32, up to 100 epochs with early stopping (patience 10), and MSE loss (binary cross-entropy for Wu).

2.9 Design Task Evaluation

The Grasso library contains designed signal peptide variants with known WA values (Set = NaN in the original data, after excluding rows that overlap with the train/test sets). A minimum of 5 variants per gene is required for per-gene metrics, yielding 4832 variants across 131 genes for evaluation. I evaluate two feature representations: the 156 physicochemical features available from the source data, and ESM2-650M embeddings (1280 dimensions) generated via mean-pooled representations from the pretrained model [Lin et al., 2023]. For each feature type, RF and NN models are trained on the Grasso training set with the best hyperparameters from Sections 2.3–2.4 and evaluated on the design variants. Wild-type WA for each gene is determined by matching the wild-type signal peptide sequence (from the “WT sequences” sheet of the source spreadsheet) against Library rows with the same amino acid sequence and taking the mean WA of matching rows. The matching is performed against the *full* Grasso library, which contains 137 genes in total (more than the 134/135 retained in the modeling train/test subset¹), so this yields WT WA values for all 137 genes. Per-gene and overall Spearman ρ are computed, along with classification accuracy for predicting whether a variant performs better or worse than wild type.

2.10 Evaluation

All models are evaluated on the same held-out test set using mean squared error (MSE), coefficient of determination (R^2), and Spearman rank correlation (ρ). MSE is the primary metric for consistency with prior work. Because a single model varies by $\sim \pm 0.05$ MSE across random seeds, all results are reported as a mean $\pm$ standard deviation over independent retrains, with 95% bootstrap confidence intervals.

3 Results

3.1 Baseline Reproduction

Using the exact parameters from Grasso et al. [2023], I obtain a test MSE of 1.193 ($R^2 = 0.750$, Spearman $\rho = 0.859$), within 2.2% of the reported value of 1.22. My random forest reproduction overfits substantially: the training MSE is 0.303 ($R^2 = 0.932$), far below the 1.75 training MSE reported by Grasso et al. [2023], indicating that my reproduction memorizes the training set more aggressively than the original despite matching test performance. Figure 2 shows predicted versus actual WA values.

¹Gene-count denominators: the test set spans 134 genes and the training set 133 (overlapping in 132); their union is 135 (one train-only gene, two test-only); the full Grasso library has 137 genes; and 131 design genes have ≥ 5 variants each.




Figure 2: Baseline reproduction using physicochemical features and the exact random forest parameters from Grasso et al. [2023]. Training and test mean-squared error; the dotted line marks Grasso et al.’s reported test MSE of 1.22. The reproduced test MSE (1.19) matches it, while the low training MSE (0.30) reflects the random-forest overfitting discussed in the text.

3.2 Scalar Model Comparison

With the baseline established, I compare all model–feature combinations. Table 1 presents results for all scalar model–feature combinations.

Table 1: Test set performance across all scalar model and feature type combinations, sorted by MSE. The Δ column shows percent change relative to the Grasso et al. [2023] baseline MSE of 1.22.

Feature Type	Model	Test MSE	Test R^2	Spearman ρ	Δ
Ginkgo AA-0	NN	1.050	0.780	0.866	−14.0%
ESM2-3B	NN	1.067	0.776	0.861	−12.5%
ESM2-650M	NN	1.090	0.771	0.861	−10.7%
Ginkgo AA-0	RF	1.158	0.757	0.857	−5.1%
ESM2-650M	RF	1.189	0.750	0.853	−2.5%
PhysChem	RF (baseline)	1.193	0.750	0.859	−2.2%
PhysChem	RF (tuned)	1.208	0.747	0.857	−1.0%
ESM2-3B	RF	1.245	0.739	0.849	+2.0%
PhysChem	NN	1.328	0.721	0.840	+8.8%

Figure 3 visualizes these results as a grouped bar chart.




Figure 3: Test MSE across feature types and model architectures. The dotted line marks the Grasso baseline MSE of 1.22. Lower values indicate better performance.

Random Forest Results. Tuned RF models achieve MSE values between 1.16 and 1.25 across all feature types. Ginkgo AA-0 embeddings yield the best RF result (MSE 1.158), followed by ESM2-650M (1.189) and physicochemical features (1.208). ESM2-3B performs worst (MSE 1.245), likely due to the curse of dimensionality in tree-based splitting with 2560 features.

Neural Network Results. The ranking changes substantially for neural networks. All three PLM embedding types outperform the baseline, with Ginkgo AA-0 achieving the best scalar result (MSE 1.050, $R^2 = 0.780$). However, the NN trained on physicochemical features performs worse than the RF baseline (MSE 1.328 vs. 1.193).

3.3 Vector Regression

Given that neural networks with PLM embeddings outperform all other scalar approaches, I next ask whether predicting the full bin probability distribution can yield further gains. Table 2 presents results for vector regression models.

Table 2: Vector regression test set performance (5-seed ensembles, $N_{\text{test}} = 1273$). WA is recovered from predicted bin probabilities. See text for caveats on comparing with Table 1.

Embedding	Loss	Test MSE	Test R^2	Spearman ρ
Ginkgo AA-0	Focal	1.001	0.790	0.867
Ginkgo AA-0	Cross-entropy	1.003	0.790	0.867
ESM2-650M	Focal	1.049	0.780	0.859
ESM2-650M	Cross-entropy	1.054	0.779	0.861
ESM2-3B	Cross-entropy	1.109	0.767	0.857
ESM2-3B	Focal	1.117	0.766	0.856

The best vector regression model (Ginkgo AA-0 + focal loss) achieves $\text{MSE} = 1.001$, a 4.7% improvement over the best scalar NN (1.050) and an 18% improvement over the baseline (1.22). Focal loss and cross-entropy perform comparably (differences of ~ 0.002 – 0.008 MSE, within the per-seed noise); I use focal for consistency with the **net4** prototype. ESM2-3B again performs worst among the three PLMs, consistent with the scalar regression results. The vector regression test set (1273 samples) is smaller than the scalar test set (1326) and uses 5-seed ensembles, so Tables 1 and 2 are not directly comparable. Figure 4 shows the comparison across embeddings and loss functions.




Figure 4: Vector regression results comparing three PLM embedding types under cross-entropy versus focal loss (5-seed ensemble test MSE).

3.4 Architecture and Regularization Optimization

The base vector model improves over scalar regression, but further gains are possible. Table 3 presents results from the full-data optimization described in Section 2.7.

Table 3: Full-data vector regression optimization ($N_{\text{train}} = 3068$, $N_{\text{test}} = 1326$, Ginkgo AA-0 + focal loss). Each row reports the 5-seed-ensemble test MSE.

Phase	Configuration	Dropout	Test MSE (mean $\pm$ std)
Dropout	(256, 256, 128)	0.20	0.978 ± 0.012
	(256, 256, 128)	0.35	0.957 ± 0.009
4-layer	(256, 256, 128, 64)	0.20	~ 0.94 – 0.95
	(384, 384, 256, 128)	0.20	~ 0.95
Mixed	3-arch $\times$ 5-seed (15 models)	—	~ 0.97

On the validation split, dropout rates in 0.20–0.35 are statistically indistinguishable

(Appendix E). Re-training the identical (256, 256, 128) configuration multiple times gives 0.978 ± 0.012 (dropout 0.20) and 0.957 ± 0.009 (dropout 0.35) on the Ginkgo AA-0 embedding. The reproducible vector-NN result on AA-0 is thus ~ 0.96 – 0.98 , a ~ 20 – 22% improvement over the published random forest baseline of 1.22 [Grasso et al., 2023]. Four-layer architectures and mixed-architecture ensembles do not improve upon the 3-layer model, and 20-seed ensembles do not improve over 5-seed.

Figure 5 shows the full optimization landscape.




Figure 5: Full-data vector regression optimization results. Configurations are grouped by optimization phase; all fall within the ~ 0.97 – 1.13 range.

3.5 A Reproducible Open-Weight Ensemble (Headline Result)

The best single embedding above is Ginkgo AA-0, but it is discontinued and cannot be regenerated by anyone (Section 2.1), so it cannot anchor a reproducible result. I therefore base the headline on *open-weight* models and ask whether an ensemble of architecturally diverse open PLMs can recover AA-0’s performance reproducibly. Generating embeddings from ProtT5 (T5 architecture) and ProtBERT (BERT architecture) alongside ESM2-650M, and ensembling the per-embedding vector NNs, gives the fully reproducible result in Table 4. These open-ensemble vector NNs use the same (256, 256, 128) layout but with PReLU activation and dropout 0.3, rather than the LeakyReLU approximation used in the analyses above: PReLU matches the `ParametricRampLayer` in Dr. Schrier’s net4 prototype, and adopting it yields a small but reproducible ~ 0.02 MSE improvement (the only single-model change that helped). The dropout of 0.3 lies within the validation-indistinguishable 0.20–0.35 range (Appendix E).

Table 4: Open-weight ensemble result (vector NN, 5-seed ensembles). The ProtT5 + ESM2-650M + ProtBERT ensemble is fully reproducible from public HuggingFace weights and recovers all but ~ 0.01 of the now-unavailable Ginkgo AA-0 performance, improving $\sim 20\%$ over the Grasso baseline. Standard deviations are reported where retrains were replicated.

Model (vector NN)	Test MSE	Reproducible?
ProtT5 + ESM2-650M + ProtBERT (open ensemble)	0.975 ± 0.003	open weights
ProtT5 + ProtBERT	0.977 ± 0.005	open weights
ProtT5 + ESM2-650M	0.994 ± 0.004	open weights
ProtT5 (alone)	1.013 ± 0.007	open weights
ProtBERT (alone)	1.068 ± 0.008	open weights
ESM2-650M (alone)	1.070 ± 0.004	open weights
Ginkgo AA-0 (alone)	0.965 ± 0.005	<i>API discontinued</i>
Grasso et al. RF (baseline)	1.22	—

The open-model ensemble (**0.975 ± 0.003**) closes the gap to the now-unavailable AA-0 (0.965 ± 0.005) to within ~ 0.01 and improves $\sim 20\%$ over the Grasso RF baseline (1.22), and *anyone can regenerate every embedding from public weights with no API key and no defunct service*. The ensemble works here because ProtT5 (1.013) is strong enough that the three open models are quality-comparable, so averaging their partially independent errors yields a genuine improvement; ensembling AA-0 with ESM2 does not help, because AA-0 (0.965) so dominates ESM2 (1.07) that averaging only dilutes the stronger model. A linear probe confirms that AA-0 and ESM2-650M capture essentially the same signal (predictions agree at Pearson $r = 0.953$, Spearman 0.952, with near-identical rank correlation to the measured WA: Spearman 0.865 vs. 0.854), so the open ensemble recovers AA-0’s information rather than merely approximating its score.




Figure 6: Reproducible open-weight headline. Test MSE for the vector NN on each open embedding and their ensembles; the discontinued, unreproducible Ginkgo AA-0 (rose) is shown only for comparison. The open ProtT5 + ESM2-650M + ProtBERT ensemble (0.975 ± 0.003) recovers AA-0 to within ~ 0.01 and improves $\sim 20\%$ over the Grasso baseline (dotted).

3.6 Ablation Analysis

To isolate the contributions of each modification, Table 5 presents a cumulative ablation starting from the validation-split baseline. All entries are 5-seed ensembles.

Table 5: Cumulative ablation of factors contributing to the reproducible ~ 0.96 Ginkgo AA-0 result. Each row adds one change. The first row uses the validation-split test set ($N = 1273$); subsequent rows use the full test set ($N = 1326$).

Configuration	5-Seed MSE	Mean 1-Seed	Δ
Val-split, (256, 256), drop=0.20	1.001	1.068	—
+ Full-data training	1.001	1.059	0.0%
+ Deeper arch (256, 256, 128)	0.993	1.039	−0.8%
+ Higher dropout (0.35)	0.957 ± 0.009	0.979	−3.6%

Full-data training alone provides little benefit at the ensemble level, though it modestly improves single-seed performance. Adding the third hidden layer and increasing dropout provide small gains that are within seed-to-seed variability (± 0.01 on the ensemble); the dominant, reliable gain is from ensembling multiple seeds, which reduces variance relative to single models. The small per-step gains fall within seed-to-seed variability, so the ablation indicates direction of effect rather than precise per-step magnitudes.

3.7 Output Formulation Comparison

To disentangle the contributions of the output formulation (10-bin vs. fewer classes) from the loss function (focal vs. CCE), I compare five output strategies on the same architecture and data (Table 6).

Table 6: Output formulation comparison. All models use the same architecture (256, 256, 128; LeakyReLU; dropout 0.35) and 5-seed ensembles on Ginkgo AA-0 embeddings. MSE is computed against continuous WA via class-centroid weighting for classification models.

Formulation	Loss	Dim	MSE	MSE 95% CI	Spearman ρ
10-bin vector	focal	10	0.971	[0.855, 1.097]	0.870
10-bin vector	CCE	10	0.994	[0.877, 1.121]	0.867
5-class	CCE	5	1.126	[1.001, 1.257]	0.857
3-class	CCE	3	1.293	[1.161, 1.433]	0.844
Binary	BCE	2	1.241	[1.111, 1.380]	0.833

The 10-bin CCE model (MSE 0.994) outperforms all classification baselines (MSE 1.13–1.29), confirming that the improvement arises from the output formulation (predicting over 10 bins rather than fewer classes) and not solely from focal loss. Within the 10-bin formulation, focal loss and cross-entropy are statistically indistinguishable (0.971 vs. 0.994, within the per-seed noise); I report focal for consistency with the `net4` prototype. (The 0.971 here

is one retrain of a configuration whose reproducible value is $\sim 0.96\text{--}0.98$ over seeds; Table 3.) MSE increases monotonically with decreasing output dimensionality from 10 to 3 classes (+33%), though binary classification (MSE 1.241) slightly outperforms 3-class (MSE 1.293), likely because the sigmoid output produces continuous interpolation between two centroids rather than discrete class assignments. Figure 7 shows the comparison.




Figure 7: Output formulation comparison: test MSE for predicting the 10-bin distribution (focal and cross-entropy losses) versus 5-, 3-, and 2-class formulations. The 10-bin vector regression is best; the gain comes from the distributional output, not the loss function.

3.8 Embedding Contribution Analysis

To quantify how much predictive power comes from the PLM embedding versus the MLP architecture, I compare the best NN against three simpler baselines using the same Ginkgo AA-0 embeddings (Table 7).

Table 7: Baseline comparison on Ginkgo AA-0 embeddings ($N_{\text{test}} = 1326$). Linear probe uses the same focal loss and 5-seed ensemble as the NN; Ridge and XGBoost predict scalar WA. The Best NN MSE is the reproducible mean for this configuration (0.957 ± 0.009 over independent retrains; Section 4); the bootstrap CI omits initialization variance and is a complementary single-draw bound.

Model	MSE	MSE 95% CI	Spearman ρ	R^2
Linear probe (Input $\rightarrow$ softmax)	1.039	[0.926, 1.161]	0.865	0.782
Ridge regression	1.088	[0.984, 1.198]	0.859	0.772
XGBoost	1.068	[0.958, 1.184]	0.862	0.776
Best NN (256, 256, 128)	0.957	[0.832, 1.066]	0.872	0.802

Relative to a mean-prediction baseline (MSE 4.76), the linear probe captures 97.9% of the total improvement achieved by the best NN, with the three hidden layers contributing only an additional 2.1%. Measured instead against the best RF (MSE 1.16; Table 1), the

linear probe captures 60% of the RF-to-NN improvement, with the MLP contributing the remaining 40%. The choice of reference baseline substantially affects this decomposition, but under either framing the embedding itself encodes the majority of task-relevant information, with the MLP head providing a further, smaller nonlinear refinement. Nevertheless, the NN consistently outperforms the linear probe (0.957 vs. 1.039, 7.9% improvement), and all four models substantially outperform the RF baseline (MSE 1.16–1.25; Table 1). To test whether the (256, 256, 128) architecture is over-parameterized, I also evaluate smaller MLPs: a (64, 32) network performs comparably to the full model (within seed noise), while a (128, 64) network achieves 0.991. This suggests that much of the architecture’s capacity is redundant given the embedding’s richness; the various configurations are statistically indistinguishable within seed noise. Figure 8 shows the baseline comparison; Figure 9 presents the cumulative ablation.




Figure 8: Baseline comparison on Ginkgo AA-0 embeddings: test MSE for a linear probe (no hidden layers), Ridge, XGBoost, and the vector NN, with the Grasso baseline marked (dotted). The linear probe captures most of the embedding’s predictive power; the MLP provides a further, smaller improvement.




Figure 9: Cumulative ablation of the changes contributing to the Ginkgo AA-0 vector result: full-data training, deeper architecture, and higher dropout (Table 5). Note: the final bar is a single-initialization value; over independent retrains this configuration gives 0.957 ± 0.009 .

3.9 Gene-Stratified Evaluation

To assess whether the model generalizes across genes rather than memorizing gene-specific patterns, I perform leave-one-gene-out (LOGO) cross-validation over all 135 genes in the combined train+test data ($N = 4341$ sequences with valid bin probabilities; the combined set contains one additional train-only gene beyond the 134 in the standard partition). For each fold, all sequences from a single gene are held out, the model is trained on the remaining genes (single seed, same architecture and hyperparameters), and predictions are evaluated on the held-out gene.

LOGO CV yields an overall MSE of 2.187 [2.075, 2.304] (95% bootstrap CI), compared to 0.964 for a comparable single-seed standard model, a $2.27\times$ increase. (LOGO uses a single seed per fold due to computational constraints: 135 folds would require 675 model trainings with 5 seeds each, so the single-seed standard MSE of 0.964 is the fair comparator; the corresponding reproducible 5-seed standard ensemble is ~ 0.96 .) This realistic cross-gene generalization estimate (MSE 2.19, Spearman 0.72) is reported alongside the standard split, not behind it: the standard test MSE is optimistic precisely because 132 of 134 genes (99.8% of test sequences) overlap train and test. The Spearman correlation under LOGO is $\rho = 0.720$ [0.704, 0.734], indicating that the model retains substantial rank-ordering ability even when predicting entirely unseen genes. Per-gene MSE varies widely (median 1.61, range 0.063–12.72, SD 1.80), with outlier genes such as *ypjP* (MSE 12.72, $n = 47$) driving the right tail of the distribution. No strong relationship between gene sample size and per-gene MSE is apparent, suggesting that prediction difficulty is driven by sequence-level factors rather than data scarcity alone.

The $2.27\times$ gap between standard and LOGO evaluation confirms that the original train/test partition, with its gene overlap, allows the model to exploit gene-specific patterns. Note that the LOGO procedure trains on the combined train+test data (~ 4300 samples per fold vs.

3068 for the standard split), which should favor the LOGO models; that the gap persists despite this advantage makes the finding more striking.

This inflation is a property of the split, not of the PLM/NN pipeline. Running the identical LOGO protocol on the Grasso random-forest baseline (156 physicochemical features, the paper’s exact RF hyperparameters) raises its MSE from 1.25 (standard split) to 2.88 (LOGO), a $2.31\times$ inflation essentially identical to the neural network’s $2.3\times$. This companion run deliberately uses the same row set for its standard and leave-one-gene-out splits (applying no length filter), which keeps the $2.31\times$ ratio apples-to-apples; its 1.25 standard-split MSE recovers the 1.19 baseline reproduction reported above once the 10–40 amino acid length filter is applied. A random forest on physicochemical features and a neural network on PLM embeddings inflate by the same factor, confirming that the gap is driven by the $\sim 99.8\%$ gene overlap in the benchmark split rather than by the model class.

To decompose this gap, I also run a standard (non-gene-stratified) random 5-fold CV on the same combined data, yielding MSE 1.226 [1.151, 1.302]. Comparing the three evaluation protocols (standard test 0.964, random CV 1.226, LOGO 2.187) reveals that 21% of the LOGO gap arises from inherent cross-validation variance (the penalty of predicting held-out data rather than data from the same distribution as training), while the remaining 79% is attributable to gene-identity leakage. This confirms that the dominant source of optimistic bias in the standard evaluation is gene overlap, not simply the use of a fixed test set.

Nevertheless, the LOGO Spearman ρ of 0.720 demonstrates meaningful cross-gene predictive ability, and the model still substantially outperforms a mean-prediction baseline (MSE 4.76) even under the stricter evaluation. The bootstrap CI on the LOGO result captures test-set sampling variance but not model initialization variance, since all 135 folds use a single fixed seed (42); the true uncertainty is therefore somewhat larger than the reported interval.




Figure 10: Gene-stratified (LOGO) evaluation: standard (gene-overlapping) test MSE (0.96, single-seed) versus LOGO CV MSE (2.19). The $\sim 2.3\times$ gap shows that the standard split (in which 132 of 134 genes overlap train and test) substantially overestimates cross-gene generalization.

3.10 Cross-Dataset Generalization

Table 8 summarizes cross-dataset performance for the best ESM2-650M RF and 5-seed NN ensemble.

Table 8: Cross-dataset generalization. Models trained on Grasso data are evaluated on four external datasets. Spearman correlations are sign-aligned so that positive indicates agreement (external assays score higher = better secretion, whereas predicted Grasso WA is lower = better; Section 2.8). All are positive; see text for multiple-testing correction.

Dataset	N	RF Spearman (p)	NN Spearman (p)
Wu	81	+0.314 (0.004)	+0.279 (0.012)
Xue	322	+0.278 (<0.001)	+0.138 (0.013)
Zhang-P43	114	+0.218 (0.020)	+0.191 (0.042)
Zhang-PglVM	114	+0.250 (0.007)	+0.211 (0.024)

All four external datasets yield small but consistently *positive* agreement-aligned correlations for both RF and NN, indicating that the secretion-quality ranking transfers weakly but in the biologically correct direction. The correlations are individually significant ($p < 0.05$); under Bonferroni correction across the eight zero-shot tests ($\alpha = 0.00625$), the Wu and Xue random-forest results survive, while the remaining (Zhang and NN) rows are directionally consistent but not individually significant after correction. The transfer is weak (~ 0.2 – 0.3 , a few percent of rank variance) and absolute magnitude does not transfer. For the Zhang datasets, the measured WA values correlate at $r = 0.977$ across the two promoters (P43 and PglVM), indicating that signal peptide rankings are largely conserved within a host. Figure 11 visualizes the cross-dataset evaluation.




Figure 11: Cross-dataset generalization results. Zero-shot Spearman correlations between predicted and observed signal peptide efficiency across four external datasets, for the random forest (RF) and neural network (NN), sign-aligned so positive indicates agreement (Section 2.8). All correlations are positive (weak transfer in the biologically correct direction); per-dataset p -values are given in Table 8.

Fine-tuning results. Given the weak but consistently positive zero-shot transfer, I next test whether fine-tuning the pretrained vector ensemble on target-domain data improves it. Table 9 compares zero-shot and fine-tuned performance of the ESM2-650M vector regression

ensemble on each external dataset; the zero-shot values differ from Table 8 because they come from the vector ensemble rather than the scalar RF/NN models.

Table 9: Cross-dataset fine-tuning results. The pretrained vector regression ensemble is fine-tuned on each external dataset using 5-fold CV $\times$ 5 seeds. Zero-shot uses the vector model directly; fine-tuned replaces and retrains the output head on the external metric. Zero-shot correlations are sign-aligned (positive = agreement, as in Table 8); the fine-tuned head is trained on that metric directly, so a negative value indicates genuine anti-correlation rather than a convention flip. Both correlations are computed as a single *pooled out-of-fold* Spearman ρ over all held-out predictions, with a 5000-resample bootstrap 95% CI; significance is judged by whether the CI excludes zero. (A mean of per-fold correlations is upward-biased in magnitude and would obscure the Wu effect, so the pooled estimator is used.)

Dataset	N	Zero-Shot ρ [95% CI]	Fine-Tuned ρ [95% CI]
Wu (binary)	81	+0.246 [+0.034, +0.455]	+0.320 [+0.101, +0.516]
Xue	322	+0.161 [+0.053, +0.264]	−0.196 [−0.300, −0.091]
Zhang-P43	114	+0.220 [+0.040, +0.377]	−0.307 [−0.468, −0.129]
Zhang-PglVM	114	+0.188 [+0.014, +0.346]	−0.355 [−0.504, −0.187]

On the aligned pooled out-of-fold estimate, the zero-shot transfer is weakly positive and significant for all four datasets (Spearman ρ from +0.16 to +0.25, each 95% CI excluding zero, marginally so for Zhang-PglVM). Fine-tuning retrains the output head on the target metric over frozen Grasso features, so the fine-tuned column reflects a target-supervised head on those features rather than pure transfer. For the binary Wu dataset, a Wu-supervised sigmoid head reaches +0.320 [+0.101, +0.516], comparable to the +0.246 zero-shot value; the two CIs overlap, so the change is not significant. For the three continuous datasets, the fine-tuned linear head *significantly degrades* the result, driving the correlation from significantly positive to significantly negative with the zero-shot and fine-tuned CIs separated in every case (Xue +0.161 $\rightarrow$ −0.196, Zhang-P43 +0.220 $\rightarrow$ −0.307, Zhang-PglVM +0.188 $\rightarrow$ −0.355). Limited target-domain data thus does not merely fail to help; it destroys the weak working transfer present zero-shot, consistent with the head overfitting spurious correlations in the small folds (64–258 samples per fold). In short, the rank ordering transfers weakly zero-shot, and a target-supervised head on frozen features does not improve it and, for the continuous assays, harms it. Figure 12 visualizes the zero-shot versus fine-tuned performance.




Figure 12: Cross-dataset fine-tuning results (sign-aligned so positive = agreement). Left: zero-shot vs. fine-tuned pooled out-of-fold Spearman correlations for the vector regression ensemble on four external datasets; error bars are 95% bootstrap confidence intervals. Right: the change from fine-tuning (fine-tuned minus zero-shot), which is positive but non-significant for the binary Wu task and significantly negative for the three continuous datasets.

3.11 Design Task Evaluation

Beyond generalization to external datasets, I also evaluate whether the models can guide the design of new signal peptide variants within the Grasso system. Table 10 summarizes model performance on ranking 4832 designed signal peptide variants across 131 genes (of 134 total, after requiring ≥ 5 variants per gene).

Table 10: Design task evaluation across two feature representations ($N = 4832$ variants across 131 genes). Models predict relative performance of designed signal peptide variants compared to wild type.

Features	Model	Spearman ρ	Pearson r	MSE	ClassAcc (%)
ESM2-650M (1280d)	NN	0.390	0.403	4.05	72.0
PhysChem (156d)	NN	0.378	0.388	4.02	72.5
PhysChem (156d)	RF	0.369	0.378	4.07	74.2
ESM2-650M (1280d)	RF	0.363	0.367	4.15	72.9

All four models achieve Spearman correlations above 0.36 and classification accuracy of 72–74%, with all correlations highly significant ($p < 10^{-100}$). The ESM2-650M NN achieves the highest rank correlation ($\rho = 0.390$), outperforming the PhysChem NN (0.378), PhysChem RF (0.369), and ESM2-650M RF (0.363). RF models achieve marginally better classification accuracy (74.2% PhysChem, 72.9% ESM2-650M) than NNs (72.0–72.5%). The improvement from ESM2-650M over PhysChem features is modest for the design task, consistent with the smaller RF–NN gap observed for scalar regression (Section 3).

To compare directly with the design validation in Grasso et al. [2023] (their Figure 4e: 11 of 15 designed peptides fell within ± 1 WA of prediction, 73%), I compute the fraction of designed variants whose predicted WA lands within ± 1 of the measured value. All four models achieve $\sim 41\%$ (PhysChem RF 41.1%, PhysChem NN 41.4%, ESM2 RF 41.3%, ESM2 NN 40.9%). This is lower than Grasso’s 73%, but the comparison is not apples-to-apples: their 15 designs were hand-picked and independently re-assayed in nanoliter reactors, whereas

I evaluate the *full* $\sim 4,800$ -variant library against the original (noisier) library-screening WA, an unselected, lower-bias evaluation. The parity plot (Figure 13) shows predicted versus measured WA per model.




Figure 13: Design-task parity (predicted vs. measured WA) for the full design library, by featurization (PhysChem = circle, ESM2 = square) and model (RF = open, NN = filled); dashed line is $y = x$. Comparable metric to Grasso Fig. 4e: $\sim 41\%$ of variants fall within ± 1 WA (vs. their $11/15 = 73\%$ on 15 curated, independently re-assayed designs).

4 Discussion

Table 11: Summary of key results.

Result	Value	Reference
Reproducible headline: open-weight ensemble MSE	0.975 ± 0.003	Table 4
Improvement over baseline (1.22)	$\sim 20\%$	Table 4
ESM2-650M single-embedding MSE (open)	1.070 ± 0.004	Table 4
Ginkgo AA-0 (best, but API discontinued)	0.965 ± 0.005	Table 4
Reproducible AA-0 vector NN (mean $\pm$ std)	0.957 ± 0.009	Table 3
Linear probe MSE (vs. NN)	1.039 (most of improvement)	Table 7
LOGO CV MSE (realistic generalization)	2.19 (Spearman 0.72)	Section 3.9
Design task Spearman ρ	0.363–0.390	Table 10

Model–Representation Interaction. The interaction between model type and feature representation (Table 1) likely reflects a fundamental difference in how RF and NN handle high-dimensional inputs. RF’s random subspace sampling provides robustness to uninformative dimensions but limits its ability to exploit complex feature interactions, explaining why it plateaus across all four feature sets. Neural networks can learn nonlinear combinations of PLM embedding dimensions that capture secretion-relevant information, but require sufficient input richness to outperform simpler methods. The poor NN performance on

physicochemical features (MSE 1.328) likely reflects a dimensionality mismatch: with only 156 input dimensions and ~ 3000 training samples, the neural network cannot learn useful representations beyond what RF already captures through axis-aligned splits, and the higher capacity becomes a liability rather than an asset.

PLM Model Size. A related question is whether larger PLMs automatically produce better predictions. Surprisingly, the largest PLM (ESM2-3B, 2560 dimensions) does not yield the best results. For RF, ESM2-3B performs worst among all models. For NN, ESM2-3B (MSE 1.067) outperforms ESM2-650M (1.090) but is still outperformed by Ginkgo AA-0 (1.050). With only ~ 3000 training samples, the additional dimensions of the 3B model may introduce noise without proportional signal. The Ginkgo AA-0 model, despite having fewer parameters than ESM2-3B, may encode more task-relevant biological information for signal peptide function.

Vector Regression. Beyond the choice of input features, the regression formulation itself also matters. Predicting bin probability distributions instead of scalar WA values yields a 4.7% MSE improvement over the best scalar NN ($1.050 \rightarrow 1.001$). This gain likely arises from the richer supervision signal: each training example provides 10 target values instead of one, and the softmax output naturally enforces a valid probability distribution. Focal loss and cross-entropy are statistically indistinguishable (Section 3.7); the gain comes from predicting 10-bin distributions rather than from the loss function, as the 10-bin CCE model (MSE 0.994) outperforms all classification baselines (MSE 1.13–1.29). However, the vector and scalar NNs also differ in activation function (LeakyReLU vs. ReLU) and the absence of batch normalization (Appendix F), so the improvement cannot be attributed solely to the distributional output without further ablation. With an 80/20 validation split, the best vector model achieves MSE 1.001, already an 18% improvement over the published baseline of 1.22. The reproducible vector-NN result on Ginkgo AA-0 is ~ 0.96 – 0.98 (e.g. 0.957 ± 0.009 at dropout 0.35), a ~ 20 – 22% improvement over the baseline (Section 3.4). The key ingredients are (i) an additional hidden layer providing greater model capacity, (ii) dropout regularization, and (iii) multi-seed ensembling to reduce prediction variance. The reproducible contribution is methodological as much as numerical: ensembling, which reduces the $\sim \pm 0.05$ per-model seed variance, is the most reliable single intervention.

Reproducibility and high variance. The per-seed variance and a cross-check against the Wolfram `net4` prototype are detailed in the repository’s reproducibility notes (`docs/reproducibility_findings.md`).

Cross-Dataset Generalization. Once the orientation conventions are aligned (Section 2.8), the cross-dataset results indicate *weak but real* transfer rather than none: all four external datasets yield small positive agreement-aligned correlations (Spearman ~ 0.2 – 0.3), and the consistency of the sign is the key evidence. A genuine failure to transfer would scatter the correlations around zero; instead all four sit systematically positive, while the within-Grasso predicted-versus-actual correlation is $+0.87$ on the same convention, the signature of a real if weak signal rather than noise. Two of the eight zero-shot tests (the Wu and Xue random

forests) survive Bonferroni correction; the remaining rows are directionally consistent but underpowered. The transfer is nonetheless weak (a few percent of rank variance), and absolute magnitude does not transfer. This likely reflects host- and cargo-specific signal patterns learned from the Grasso data that only partially generalize, since signal peptide efficiency depends on the cargo protein, promoter, host organism, and growth conditions. A signal peptide optimized for α -amylase secretion in *B. subtilis* may be poorly suited for xylanase production in the same organism, let alone for enzyme secretion in *E. coli* or *S. cerevisiae*. The Zhang cross-promoter analysis provides complementary evidence: the same 114 signal peptides correlate at $r = 0.977$ across two promoters, confirming that signal peptide rankings are largely conserved within a system, and only weakly across systems (Table 8). In practice, the cross-system transfer is too weak to rely on, so models for signal peptide optimization still benefit from task-specific training data. Fine-tuning on target-domain data does not help: a target-supervised head on frozen Grasso features leaves the binary Wu task statistically unchanged ($+0.246 \rightarrow +0.320$, overlapping CIs) and *significantly degrades* the three continuous datasets, driving their weak positive transfer to negative with separated CIs (Xue $+0.161 \rightarrow -0.196$; Zhang-P43 $+0.220 \rightarrow -0.307$; Zhang-PglVM $+0.188 \rightarrow -0.355$). Limited target data thus actively harms a working signal rather than improving it. These results underscore that transfer learning for signal peptide efficiency remains severely data-limited, and that fine-tuning on small datasets risks overfitting rather than recovering transferable patterns.

Design Task. Despite the limited cross-dataset transferability, the models still prove useful within the Grasso experimental system. The design task demonstrates that models trained on characterized variants can meaningfully rank novel designed signal peptides, with Spearman $\rho = 0.36$ – 0.39 and classification accuracy of 72–74%, well above the 50% random baseline. Generating ESM2-650M embeddings for the design variants and evaluating all four model–feature combinations reveals that the ESM2-650M NN achieves the highest rank correlation ($\rho = 0.390$), though the improvement over PhysChem features is modest (0.390 vs. 0.378 for NN, 0.363 vs. 0.369 for RF). This contrasts with the larger PLM advantage observed in the scalar regression task, suggesting that the design task’s evaluation structure (per-gene ranking rather than absolute WA prediction) may limit the benefit of richer representations. The moderate but significant correlations suggest that these models can serve as a practical filter for prioritizing candidates in experimental screening campaigns.

Limitations. Several methodological caveats apply. The RF and NN hyperparameter searches used different validation strategies (5-fold CV vs. single train/val split), which may introduce some asymmetry in the comparison. The full-data optimized models train without a validation split, which precludes early stopping; I compensate with dropout and ensembling. Appendix E shows that all dropout values ≥ 0.20 produce statistically indistinguishable validation MSE. The external fine-tuning datasets are small (81–322 sequences), so even the pooled out-of-fold correlations I report carry wide bootstrap intervals; the cross-dataset conclusions should be read as directional rather than precise. Validation MSE improves from 1.375 at dropout 0.15 to 1.333–1.347 for all dropout rates ≥ 0.20 , with no statistically significant differences among them (all within seed-to-seed variability of ± 0.07 – 0.09); the test-

optimal dropout (0.35) produces a validation MSE of 1.347 ± 0.072 , indistinguishable from the validation-optimal (0.20, MSE 1.333 ± 0.084). The differences between top-performing models are modest; 95% bootstrap confidence intervals for all models are reported in Appendix D.

In terms of scope, the dataset is limited to *B. subtilis* with a single reporter enzyme. Ginkgo AA-0 embeddings yield the best single-embedding result but are derived from a proprietary, now-discontinued model and are permanently unreproducible; the reproducible open-weight ensemble (ProtT5 + ESM2-650M + ProtBERT, 0.975 ± 0.003) recovers all but ~ 0.01 of that performance, confirming that the central result is not contingent on proprietary representations. Cross-dataset and fine-tuning evaluations use ESM2-650M embeddings (open and reproducible), so the transferability of the Ginkgo AA-0 representations specifically remains untested. The vector regression architecture uses LeakyReLU activation, omits batch normalization, and does not apply L2 kernel regularization, while the scalar NN uses ReLU with batch normalization and L2 regularization; a systematic ablation of these architectural differences was not performed, so the vector regression improvement cannot be attributed solely to the distributional output. The original Grasso train/test partition contains overlapping genes, meaning the model may partly learn gene-specific patterns rather than generalizable signal peptide features. A LOGO cross-validation (Section 3.9) addresses this concern.

5 Conclusion

Across physicochemical descriptors and several PLM embeddings, a vector regression neural network with multi-seed ensembling substantially improves over the random forest baseline. The reproducible headline is an ensemble of architecturally diverse *open-weight* PLMs (ProtT5, ESM2-650M, ProtBERT) at MSE 0.975 ± 0.003 , a $\sim 20\%$ improvement over the 1.22 baseline and regenerable by anyone from public weights; the discontinued proprietary Ginkgo AA-0 (0.965 ± 0.005) is marginally stronger but reported only as a caveated comparison.

Two results temper this. A linear probe on the embeddings captures most of the predictive gain, so the embedding carries most of the task-relevant signal; and LOGO cross-validation (MSE 2.19, Spearman 0.720) shows that the standard split substantially overstates cross-gene generalization, so gene-stratified evaluation should be the default. Across organisms the models transfer only weakly zero-shot, in the biologically correct direction (agreement-aligned Spearman ~ 0.2 – 0.3), and small-sample fine-tuning does not improve this, but within the training domain they rank designed variants usefully (Spearman up to 0.390, 72–74% classification accuracy).

Future work could explore end-to-end fine-tuning of PLM weights on secretion-efficiency data, region-aware embedding pooling, or extension to additional host organisms and cargo proteins.

Data Availability

The original dataset is available as supplementary material to Grasso et al. [2023]. PLM embeddings were generated using publicly available model weights. All analysis code and

results are available at <https://github.com/Mehak-W/signal-peptide-secretion>.

References

- Stefano Grasso, Valentina Dabene, Margriet M. W. B. Hendriks, Priscilla Zwartjens, René Pellaux, Martin Held, Sven Panke, Jan Maarten van Dijl, Andreas Meyer, and Tjeerd van Rij. Signal peptide efficiency: From high-throughput data to prediction and explanation. *ACS Synthetic Biology*, 12(2):390–404, 2023. doi: 10.1021/acssynbio.2c00328.
- Zeming Lin, Halil Akin, Roshan Rao, Brian Hie, Zhongkai Zhu, Wenting Lu, Nikita Smetanin, Robert Verkuil, Ori Kabeli, Yaniv Shmueli, Allan dos Santos Costa, Maryam Fazel-Zarandi, Tom Sercu, Salvatore Candido, and Alexander Rives. Evolutionary-scale prediction of atomic-level protein structure with a language model. *Science*, 379(6637):1123–1130, 2023. doi: 10.1126/science.ade2574.
- Felix Teufel, José Juan Almagro Armenteros, Alexander Rosenberg Johansen, Magnús Halldór Gíslason, Silas Irby Pihl, Konstantinos D. Tsirigos, Ole Winther, Søren Brunak, Gunnar von Heijne, and Henrik Nielsen. SignalP 6.0 predicts all five types of signal peptides using protein language models. *Nature Biotechnology*, 40(7):1023–1025, 2022. doi: 10.1038/s41587-021-01156-3.
- Castrense Savojardo, Pier Luigi Martelli, Piero Fariselli, and Rita Casadio. DeepSig: Deep learning improves signal peptide detection in proteins. *Bioinformatics*, 34(10):1690–1696, 2018. doi: 10.1093/bioinformatics/btx818.
- Roshan Rao, Nicholas Bhatt, Nikunj Saunshi, Scott Lundberg, Robert Verkuil, and Alexander Rives. Evaluating protein transfer learning with TAPE. In *Advances in Neural Information Processing Systems*, volume 32, 2019.
- Ahmed Elnaggar, Michael Heinzinger, Christian Dallago, Ghalia Rehawi, Yu Wang, Llion Jones, Tom Gibbs, Tamas Feher, Christoph Angerer, Martin Steinegger, Debsindhu Bhowmik, and Burkhard Rost. ProtTrans: Toward understanding the language of life through self-supervised learning. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 44(10):7112–7127, 2022. doi: 10.1109/TPAMI.2021.3095381.
- Joshua Meier, Roshan Rao, Robert Verkuil, Jason Liu, Tom Sercu, and Alexander Rives. Language models enable zero-shot prediction of the effects of mutations on protein function. In *Advances in Neural Information Processing Systems*, volume 34, pages 29287–29303, 2021.
- Ginkgo Bioworks. AA0: A protein foundation model. Proprietary embeddings via API at <https://models.ginkgobioworks.com> (since discontinued; model weights not publicly released), 2024. Precomputed embeddings used in this study.
- Tsung-Yi Lin, Priya Goyal, Ross Girshick, Kaiming He, and Piotr Dollár. Focal loss for dense object detection. In *Proceedings of the IEEE International Conference on Computer Vision*, pages 2980–2988, 2017. doi: 10.1109/ICCV.2017.324.

Zachary Wu, Kevin K. Yang, Michael J. Liszka, Alycia Lee, Alina Batzilla, David Wernick, David P. Weiner, and Frances H. Arnold. Signal peptides generated by attention-based neural networks. *ACS Synthetic Biology*, 9(8):2154–2161, 2020. doi: 10.1021/acssynbio.0c00219.

Songlyu Xue, Xiufang Liu, Yuyang Pan, Chufan Xiao, Yunzi Feng, Lin Zheng, Mouming Zhao, and Mingtao Huang. Comprehensive analysis of signal peptides in *Saccharomyces cerevisiae* reveals features for efficient secretion. *Advanced Science*, 10(2):2203433, 2023. doi: 10.1002/advs.202203433.

Weiwei Zhang, Mingming Yang, Yuedong Yang, Jian Zhan, Yaoqi Zhou, and Xin Zhao. Optimal secretion of alkali-tolerant xylanase in *Bacillus subtilis* by signal peptide screening. *Applied Microbiology and Biotechnology*, 100(20):8745–8756, 2016. doi: 10.1007/s00253-016-7615-4.

A Hyperparameter Search Results

Tables 12 and 13 report the top-5 configurations per feature type from the randomized hyperparameter searches (100 RF configs with 5-fold CV; 40 NN configs with 80/20 validation split). Full trial data are available in the GitHub repository.

Table 12: Top-5 random forest configurations per feature type, ranked by 5-fold cross-validation MSE (100 configs searched per type).

Feature	Rank	n_{est}	Depth	Split	Leaf	CV MSE
PhysChem	1	75	None	5	2	2.019
	2	50	25	2	1e-4	2.021
	3	150	30	2	1e-4	2.028
	4	100	30	10	1e-3	2.029
	5	150	25	2	2	2.029
ESM2-650M	1	300	25	1e-3	4	1.843
	2	200	20	2	1e-3	1.845
	3	300	None	10	4	1.845
	4	75	25	5	1e-3	1.847
	5	200	20	2	4	1.847
ESM2-3B	1	150	25	2	2	1.925
	2	300	25	1e-3	4	1.927
	3	100	30	10	1e-3	1.928
	4	200	25	2	1e-3	1.929
	5	200	25	1e-3	1e-3	1.929
Ginkgo AA-0	1	300	None	10	4	1.723
	2	75	None	2	1e-4	1.723
	3	150	30	2	1e-4	1.724
	4	200	20	2	1e-3	1.724
	5	100	25	2	2	1.725

Table 13: Top-5 neural network configurations per feature type (non-log-transformed targets), ranked by validation MSE (40 configs searched per type, 80/20 train/val split). Log-transformed target variants were also evaluated during the search but are excluded here since the optimal skewness was below the transformation threshold (Section 2.2). Full trial data are available in the GitHub repository.

Feature	Rank	Layers	Drop	L2	LR / Batch	Val MSE
PhysChem	1	256, 128	0.3	0.01	1e-3 / 64	1.390
	2	128, 64	0.3	0.01	1e-3 / 32	1.425
	3	128, 64	0.4	1e-4	1e-3 / 64	1.515
	4	128, 64	0.5	0.01	5e-4 / 64	1.529
	5	128, 64	0.4	1e-3	5e-4 / 64	1.551
ESM2-650M	1	256, 128	0.4	0.01	1e-3 / 64	1.241
	2	256	0.3	1e-3	5e-4 / 64	1.266
	3	512, 256	0.5	0.01	5e-4 / 64	1.300
	4	256, 128	0.5	1e-4	5e-4 / 64	1.331
	5	256, 128	0.4	0.01	5e-4 / 32	1.353
ESM2-3B	1	1024, 512	0.4	0.01	1e-3 / 64	1.302
	2	512, 256	0.5	1e-3	1e-3 / 32	1.316
	3	512, 256	0.4	1e-3	5e-4 / 32	1.333
	4	512, 256	0.3	1e-3	1e-3 / 64	1.377
	5	512, 256	0.2	0.01	5e-4 / 32	1.387
Ginkgo AA-0	1	256, 128	0.3	1e-4	5e-4 / 32	1.146
	2	512, 256	0.5	1e-3	5e-4 / 64	1.197
	3	256	0.2	1e-3	1e-3 / 32	1.224
	4	256, 128	0.5	1e-4	1e-3 / 64	1.236
	5	256, 128	0.5	0.01	1e-3 / 32	1.248

B Vector Regression Per-Seed Results

Table 14 reports the per-seed test MSE for all six embedding–loss combinations, along with the ensemble mean.

Table 14: Per-seed test MSE for vector regression models. Each model uses the base architecture (2×256 hidden units, LeakyReLU, 20% dropout, softmax output) with 5 random seeds for train/validation splitting.

Embedding	Loss	Seed 42	Seed 123	Seed 456	Seed 789	Seed 1024	Ensemble
Ginkgo AA-0	Focal	1.026	1.120	1.023	1.128	1.046	1.001
Ginkgo AA-0	CE	1.053	1.102	1.054	1.111	1.050	1.003
ESM2-650M	Focal	1.084	1.158	1.100	1.169	1.125	1.049
ESM2-650M	CE	1.084	1.151	1.093	1.194	1.132	1.054
ESM2-3B	CE	1.235	1.224	1.179	1.196	1.214	1.109
ESM2-3B	Focal	1.215	1.220	1.185	1.226	1.193	1.117

C Architecture Search Landscape

Figure 14 shows the full architecture search landscape from the initial full-data vector regression exploration, which identified the (256, 256, 128) three-layer architecture as the top performer prior to the dropout and ensemble optimization in Section 3.4.




Figure 14: Architecture search results for full-data vector regression (Ginkgo AA-0 + focal loss). Configurations are ranked by test MSE on all 1326 test samples. The three-layer (256, 256, 128) architecture with 5-seed ensemble achieved the lowest MSE (0.973), motivating the subsequent dropout optimization (Section 3.4).

D Bootstrap Confidence Intervals

To quantify uncertainty in test set performance, I compute 95% bootstrap confidence intervals (10,000 resamples) for all 16 models (including the full-data optimized model). Each model is retrained with its best configuration and evaluated on the held-out test set; the test

set is then resampled with replacement to obtain the bootstrap distribution. These CIs capture uncertainty due to finite test-set size (sampling variability) but *not* model uncertainty from random initialization; for ensemble models, the ensemble prediction is fixed before bootstrapping begins. The dominant uncertainty for any single configuration is initialization variance, which these bootstrap CIs do not capture. The mean $\pm$ standard deviation over independent retrains reported in the main text is therefore the primary reliability indicator, with these bootstrap CIs as a complementary lower bound on uncertainty.

Table 15: Bootstrap 95% confidence intervals for test MSE, R^2 , and Spearman ρ across all models (10,000 bootstrap resamples). [†]Full-data model: (256, 256, 128) architecture, dropout 0.35, trained on all 3068 samples without validation split, evaluated on all 1326 test samples. *Note:* The full-data entry ([†]) is a single un-replicated draw; over independent retrains the identical configuration gives 0.957 ± 0.009 (Section 4), and the bootstrap CI (which omits initialization variance) understates its uncertainty. Point estimates may differ from the main tables because the bootstrap procedure retrains each model with a single random initialization.

Model	MSE	95% CI	R^2	95% CI	Spearman	95% CI
<i>Random Forest</i>						
RF PhysChem (base)	1.193	[1.078, 1.317]	0.750	[0.724, 0.773]	0.859	[0.844, 0.872]
RF PhysChem (tuned)	1.208	[1.094, 1.328]	0.747	[0.721, 0.770]	0.857	[0.843, 0.870]
RF ESM2-650M	1.189	[1.076, 1.310]	0.750	[0.725, 0.773]	0.853	[0.837, 0.867]
RF ESM2-3B	1.245	[1.126, 1.366]	0.739	[0.713, 0.763]	0.849	[0.833, 0.862]
RF Ginkgo AA-0	1.158	[1.043, 1.281]	0.757	[0.731, 0.781]	0.857	[0.841, 0.870]
<i>Neural Network (scalar)</i>						
NN PhysChem	1.413	[1.253, 1.582]	0.703	[0.666, 0.738]	0.828	[0.809, 0.846]
NN ESM2-650M	1.109	[0.988, 1.240]	0.767	[0.738, 0.794]	0.855	[0.840, 0.869]
NN ESM2-3B	1.175	[1.051, 1.302]	0.753	[0.725, 0.780]	0.850	[0.835, 0.864]
NN Ginkgo AA-0	1.191	[1.076, 1.317]	0.750	[0.723, 0.774]	0.848	[0.833, 0.861]
<i>Vector regression (5-seed ensemble)</i>						
Vec ESM2-650M CE	1.051	[0.937, 1.175]	0.780	[0.753, 0.804]	0.861	[0.845, 0.875]
Vec ESM2-650M focal	1.045	[0.934, 1.167]	0.781	[0.755, 0.804]	0.860	[0.844, 0.874]
Vec ESM2-3B CE	1.122	[1.007, 1.247]	0.765	[0.737, 0.789]	0.857	[0.841, 0.870]
Vec ESM2-3B focal	1.091	[0.981, 1.211]	0.771	[0.745, 0.794]	0.859	[0.844, 0.872]
Vec Ginkgo AA-0 CE	0.995	[0.887, 1.112]	0.791	[0.766, 0.814]	0.867	[0.852, 0.880]
Vec Ginkgo AA-0 focal	0.997	[0.887, 1.115]	0.791	[0.765, 0.814]	0.866	[0.850, 0.880]
<i>Full-data optimized (5-seed ensemble, no val split)</i>						
Vec Ginkgo AA-0 focal [†]	0.932	[0.823, 1.054]	0.804	[0.779, 0.828]	0.873	[0.857, 0.887]

Figure 15 presents these intervals as a forest plot, showing the relative performance and uncertainty of all 16 models.




Figure 15: Forest plot of test MSE with 95% bootstrap confidence intervals for all 16 models. Light-gray markers indicate scalar models (RF and NN); dark-gray markers indicate vector regression ensembles. The dotted line marks the Grasso physicochemical baseline (MSE 1.22). These CIs omit initialization variance and so understate uncertainty for the high-variance vector models (Section 4).

E Dropout Validation

To confirm that the dropout selection in Section 3.4 is not an artifact of test-set model selection, I repeated the dropout sweep using an independent 80/20 train/validation split from the training data. For each dropout rate in $\{0.15, 0.20, 0.25, 0.30, 0.35\}$, I trained 5-seed ensembles on 80% of the training data and evaluated on the held-out 20% validation set. The same (256, 256, 128) architecture, focal loss, and training protocol from Section 3.4 were used.




Figure 16: Dropout validation: MSE on held-out validation data (80/20 split) across dropout rates. The validation curve is essentially flat for dropout ≥ 0.20 , so the test-selected 0.35 and the validation-optimal 0.20 are statistically indistinguishable. Error bars represent standard deviation across 5 seeds. Absolute validation MSE is higher than test MSE because models train on only 80% of the data.

As shown in Figure 16, validation MSE drops from 1.375 at dropout 0.15 to a minimum of 1.333 at dropout 0.20, then remains between 1.338 and 1.352 for dropout 0.25–0.35; all values at dropout ≥ 0.20 are within seed-to-seed variability (± 0.07 – 0.09). The test-optimal dropout (0.35, validation MSE 1.347 ± 0.072) and the validation-optimal (0.20, 1.333 ± 0.084) are statistically indistinguishable, as are all intermediate values. This supports the conclusion that dropout values in the range 0.20–0.35 produce equivalent generalization performance, and the test-optimal choice of 0.35 is not contradicted by the validation data. The higher absolute MSE on validation (~ 1.35 vs. ~ 0.93 on test) is expected because validation models train on only 80% of the data (2454 vs. 3068 samples).

F ReLU-Squared Activation Comparison

I compared ReLU-squared ($\text{ReLU}^2(x) = \max(0, x)^2$) against the default LeakyReLU activation in the best vector regression architecture (256, 256, 128; dropout 0.35; focal loss; 5-seed ensemble). LeakyReLU achieves MSE 0.961 [0.847, 1.085], while ReLU-squared achieves MSE 5.11 [4.78, 5.44], a catastrophic degradation ($R^2 = -0.073$) worse than a mean-prediction baseline (MSE 4.76). The large per-seed variance (4.67–9.07) indicates training instability rather than a systematic accuracy–sparsity tradeoff. ReLU-squared does produce substantially sparser internal representations (Hoyer sparsity 0.67–0.92 and 56–75% exact

zeros vs. 0.27–0.32 and 0% zeros for LeakyReLU), but this comes at the cost of model collapse rather than a graceful tradeoff. This experiment uses only two learning rates (5×10^{-4} with fallback to 2×10^{-4}) and does not include BatchNorm, gradient clipping, or standard ReLU controls; the catastrophic failure may therefore reflect a hyperparameter mismatch rather than a fundamental limitation of ReLU-squared. Figure 17 shows the comparison.




Figure 17: Test MSE for LeakyReLU versus ReLU-squared activation in the vector regression architecture (256, 256, 128; dropout 0.35; focal loss; 5-seed ensemble). ReLU-squared catastrophically degrades accuracy (MSE $0.961 \rightarrow 5.11$) despite producing 2–3 $\times$ sparser internal representations (Hoyer sparsity, discussed in text); error bars are standard deviation across 5 seeds.